

DS2_midterm

My An Huynh, Soomin You

2025-03-30

```
library(tidyverse)
library(MASS)
library(caret)
library(corrplot)
library(mgcv)
library(patchwork)
library(ggplot2)
library(gtsummary)
```

Load Data

```
load("data/dat1.RData")
load("data/dat2.RData")
```

Summary Statistics

```
#train data
train = dat1 |>
  dplyr::select(-id)

x = model.matrix(log_antibody ~ ., data = train)[, -1]
y = train$log_antibody

#test data
test = dat2 |>
  dplyr::select(-id)

x_test = model.matrix(log_antibody ~ ., data = test)[, -1]
y_test = test$log_antibody

tbl_summary(train)
```

Both the train data and the test data were loaded and separated into response variable and predictor variable matrices. A summary statistics table was made to summarize the mean and proportions of the train data.

Characteristic	N = 5,000 [†]
age	60.0 (57.0, 63.0)
gender	2,427 (49%)
race	
1	3,221 (64%)
2	278 (5.6%)
3	1,036 (21%)
4	465 (9.3%)
smoking	
0	3,010 (60%)
1	1,504 (30%)
2	486 (9.7%)
height	170.1 (166.1, 174.3)
weight	80 (75, 85)
bmi	27.60 (25.80, 29.50)
diabetes	772 (15%)
hypertension	2,298 (46%)
SBP	130 (124, 135)
LDL	110 (96, 124)
time	106 (76, 138)
log_antibody	10.09 (9.68, 10.48)

[†]Median (Q1, Q3); n (%)

The following is a brief explanation of the predictor variables and the response variable summarized in the table above.

- race (1 = White, 2 = Asian, 3 = Black, 4 = Hispanic)
- gender (0 = Female, 1 = Male)
- smoking levels (0 = non-smoker, 1 = former smoker, 2 = current smoker)
- hypertension (0 = no hypertension, 1 = hypertension)
- diabetes (0 = non-diabetic, 1 = diabetic)
- SBP (systolic blood pressure in mmHg)
- LDL (LDL cholesterol in mg/dL)
- age (years)
- height (cm)
- weight (kg)
- BMI (kg/m²)
- time (time passed since vaccination in days).
- log_antibody (log transformed antibody level)

Exploratory Data Analysis

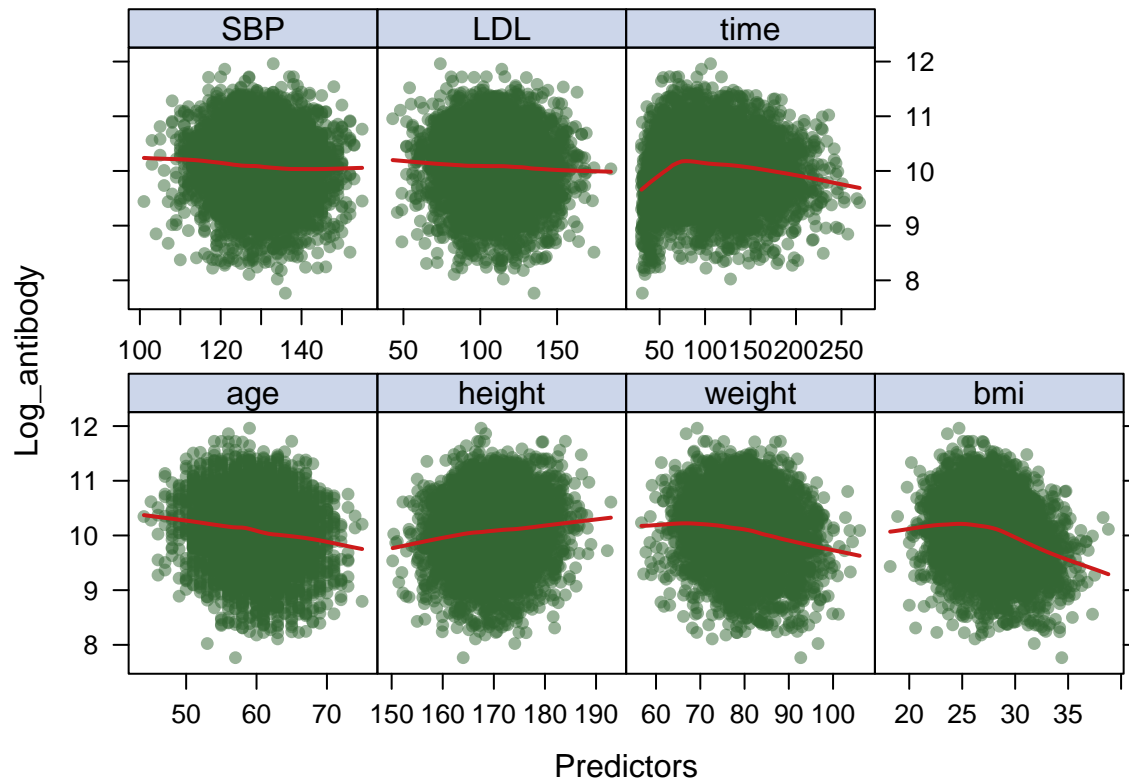
```
theme1 = trellis.par.get()
theme1$plot.symbol$col = rgb(.2, .4, .2, .5)
theme1$plot.symbol$pch = 16
theme1$plot.line$col = rgb(.8, .1, .1, 1)
theme1$plot.line$lwd = 2
```

```

theme1$strip.background$col = rgb(.0, .2, .6, .2)
trellis.par.set(theme1)

featurePlot(x = train[, -c(2,3,4,8,9,13)],
            y = train$log_antibody,
            plot = "scatter",
            span = 0.5,
            type = c("p", "smooth"),
            labels = c("Predictors", "Log_antibody"))

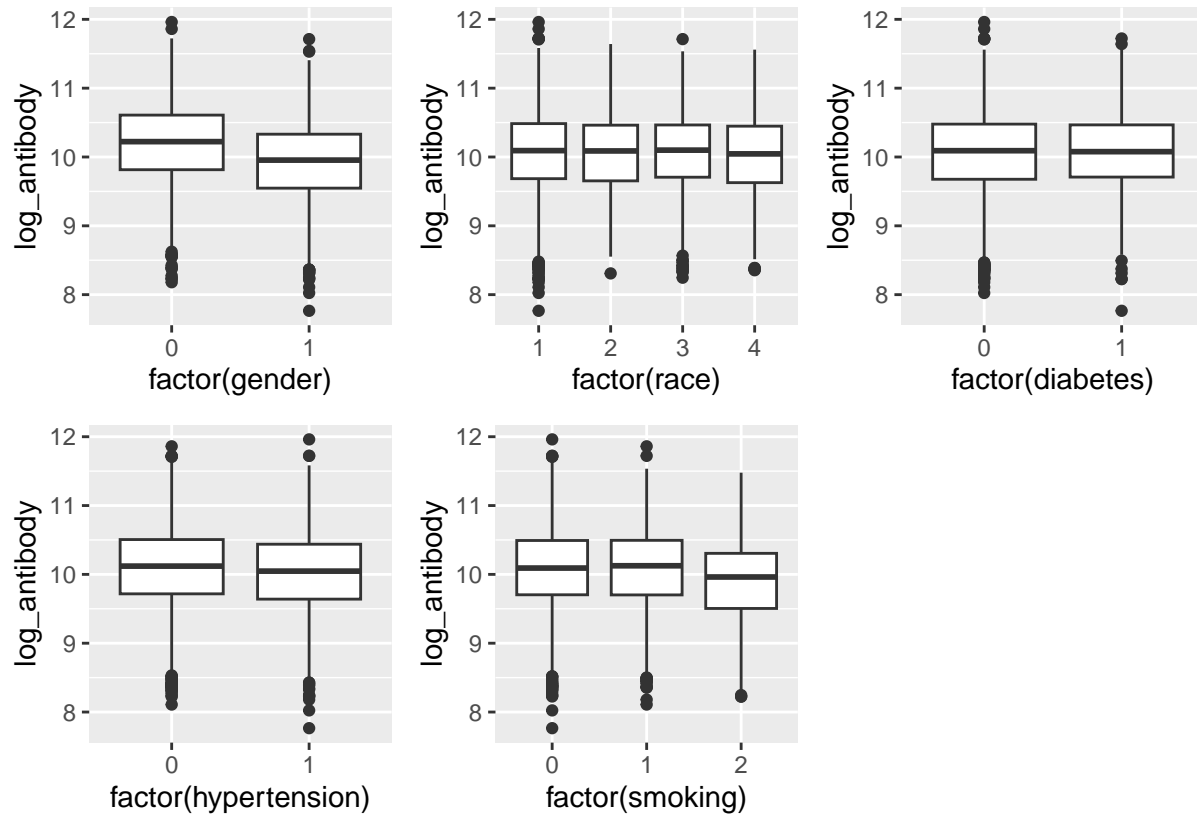
```




```

# boxplot for factor/binary variables
p1 = ggplot(aes(x = factor(gender), y = log_antibody), data = train) +
  geom_boxplot()
p2 = ggplot(aes(x = factor(race), y = log_antibody), data = train) +
  geom_boxplot()
p3 = ggplot(aes(x = factor(diabetes), y = log_antibody), data = train) +
  geom_boxplot()
p4 = ggplot(aes(x = factor(hypertension), y = log_antibody), data = train) +
  geom_boxplot()
p5 = ggplot(aes(x = factor(smoking), y = log_antibody), data = train) +
  geom_boxplot()
p1 + p2 + p3 + p4 + p5

```



```
corrplot(cor(x), method = "circle", type = "full")
```



Scatterplots were made for the numeric predictor variables and boxplots were made for the factor or binary variables.

According to the scatterplots for Systolic blood pressure (SBP), LDL cholesterol (LDL), Time since vaccination (time), Age (age), Height (height), Weight (weight), BMI (bmi), they all indicate relatively linear relationship with the response variable Log-transformed antibody level (log_antibody). Also, the boxplots indicate that the the factor and binary predictors show Gaussian characteristic.

Correlation plot was also checked and no strong multicollinearity was observed in the data.

Model training

To find an optimal prediction model of antibody levels that show the impact of the demographic and clinical factors, we decided to build various models and compare. We used linear, elastic net, PLS, MARS and GAM models as shown below.

```
#linear
ctrl1 = trainControl(method = "cv", number = 10)

set.seed(1)
linear_mod = train(x, y,
                   method = "lm",
                   trControl = ctrl1)
summary(linear_mod)

##
## Call:
## lm(formula = .outcome ~ ., data = dat)
##
```

```
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.14396 -0.35840  0.02944  0.37802  1.65090
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  26.6751961  2.3149812  11.523 < 2e-16 ***
## age         -0.0205979  0.0019385 -10.626 < 2e-16 ***
## gender      -0.2974929  0.0155977 -19.073 < 2e-16 ***
## race2       -0.0060422  0.0344613  -0.175  0.8608
## race3       -0.0075295  0.0196815  -0.383  0.7021
## race4       -0.0417571  0.0273309  -1.528  0.1266
## smoking1     0.0219907  0.0173992   1.264  0.2063
## smoking2    -0.1934834  0.0269576  -7.177 8.15e-13 ***
## height      -0.0821381  0.0135622  -6.056 1.49e-09 ***
## weight       0.0859034  0.0143481   5.987 2.29e-09 ***
## bmi         -0.2977935  0.0412612  -7.217 6.10e-13 ***
## diabetes     0.0112795  0.0215643   0.523  0.6010
## hypertension -0.0179106  0.0260931  -0.686  0.4925
## SBP          0.0015181  0.0017049   0.890  0.3733
## LDL         -0.0001645  0.0004028  -0.409  0.6829
## time        -0.0003011  0.0001795  -1.677  0.0936 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5503 on 4984 degrees of freedom
## Multiple R-squared:  0.1513, Adjusted R-squared:  0.1488
## F-statistic: 59.25 on 15 and 4984 DF, p-value: < 2.2e-16
```

```
# test error
lm_pred = predict(linear_mod, newdata = x_test)
lm_rmse = sqrt(mean((y_test - lm_pred)^2))

cat("Test error for the linear model: ", lm_rmse, "\n")
```

```
## Test error for the linear model: 0.568318
```

```
# elastic net
set.seed(1)
enet_mod = train(x = x,
                 y = y,
                 method = "glmnet",
                 tuneGrid = expand.grid(alpha = seq(0, 1, length = 25),
                                       lambda = exp(seq(-2, 2, length = 100))),
                 trControl = ctrl1)
```

```
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo,
## : There were missing values in resampled performance measures.
```

```
best_alpha = enet_mod$bestTune$alpha
best_lambda = enet_mod$bestTune$lambda
```

```
# test error
enet_pred = predict(enet_mod, newdata = x_test)
enet_rmse = sqrt(mean((y_test - enet_pred)^2))
```

```

cat("The optimal alpha for the elastic net model: ", best_alpha, "\n")

## The optimal alpha for the elastic net model: 0
cat("The optimal alpha for the elastic net model: ", best_lambda, "\n")

## The optimal alpha for the elastic net model: 0.1353353
cat("Test error for the linear model: ", enet_rmse, "\n")

## Test error for the linear model: 0.5728115

# pls model
set.seed(1)
pls_mod = train(x = x,
                y = y,
                method = "pls",
                tuneGrid = data.frame(ncomp = 1:15),
                trControl = ctrl1,
                preProcess = c("center", "scale"))
summary(pls_mod)

## Data:      X dimension: 5000 15
## Y dimension: 5000 1
## Fit method: oscorespls
## Number of components considered: 9
## TRAINING: % variance explained
##           1 comps 2 comps 3 comps 4 comps 5 comps 6 comps 7 comps
## X           10.77  19.88   31.77   38.37   43.94   47.54   53.19
## .outcome     12.83  14.29   14.38   14.41   14.42   14.45   14.51
##           8 comps 9 comps
## X           54.88  60.39
## .outcome     15.03  15.13


# test error
pls_pred = predict(pls_mod, newdata = x_test)
pls_rmse = sqrt(mean((y_test - pls_pred)^2))

cat("Test error for the linear model: ", pls_rmse, "\n")

## Test error for the linear model: 0.5690832

ggplot(pls_mod, highlight = TRUE)

```



```
# mars
set.seed(1)
mars_mod = train(x = x,
                 y = y,
                 method = "earth",
                 trControl = ctrl1)
```

```
## Loading required package: earth
## Loading required package: Formula
## Loading required package: plotmo
## Loading required package: plotrix
```

```
summary(mars_mod)
```

```
## Call: earth(x=matrix[5000,15], y=c(10.65,9.889,1...), keepxy=TRUE, degree=1,
##          nprune=9)
##
##          coefficients
## (Intercept)  10.8474469
## gender      -0.2962905
## smoking2    -0.2051269
## h(59-age)    0.0161385
## h(age-59)   -0.0229576
## h(bmi-23.7) -0.0843802
## h(27.8-bmi) -0.0619974
## h(57-time)  -0.0335293
## h(time-57)  -0.0022542
```



```
##
## Selected 9 of 10 terms, and 5 of 15 predictors (nprune=9)
## Termination condition: RSq changed by less than 0.001 at 10 terms
## Importance: bmi, gender, time, age, smoking2, race2-unused, race3-unused, ...
## Number of terms at each degree of interaction: 1 8 (additive model)
## GCV 0.2787787    RSS 1384.431    GRSq 0.2166152    RSq 0.2216218
```


```
# test error
```

```
mars_pred = predict(mars_mod, newdata = x_test)
mars_rmse = sqrt(mean((y_test - mars_pred)^2))
```

```
cat("Test error for the linear model: ", mars_rmse, "\n")
```

```
## Test error for the linear model: 0.5327718
```

```
pdp::partial(mars_mod, pred.var = c("bmi"), grid.resolution = 10) |>
  autoplot()
```



To better understand the relationship, partial dependence plot (PDP) for BMI predictor variable and the response variable was plotted.

```
# gam
set.seed(1)
gam_mod = train(x = x,
                 y = y,
                 method = "gam",
                 trControl = ctrl1)

gam_mod$bestTune
```

```

## select method
## 2 TRUE GCV.Cp
gam_mod$finalModel

##
## Family: gaussian
## Link function: identity
##
## Formula:
## .outcome ~ gender + race2 + race3 + race4 + smoking1 + smoking2 +
## diabetes + hypertension + s(age) + s(SBP) + s(LDL) + s(bmi) +
## s(time) + s(height) + s(weight)
##
## Estimated degrees of freedom:
## 0.991 0.000 0.000 4.179 7.892 1.234 0.000
## total = 23.3
##
## GCV score: 0.2786734
summary(gam_mod)

##
## Family: gaussian
## Link function: identity
##
## Formula:
## .outcome ~ gender + race2 + race3 + race4 + smoking1 + smoking2 +
## diabetes + hypertension + s(age) + s(SBP) + s(LDL) + s(bmi) +
## s(time) + s(height) + s(weight)
##
## Parametric coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 10.228177  0.015328 667.269 < 2e-16 ***
## gender      -0.297837  0.014933 -19.945 < 2e-16 ***
## race2       -0.003296  0.033009  -0.100  0.920
## race3       -0.010509  0.018837  -0.558  0.577
## race4       -0.037424  0.026176  -1.430  0.153
## smoking1     0.022219  0.016660   1.334  0.182
## smoking2    -0.193175  0.025834  -7.478 8.9e-14 ***
## diabetes     0.014230  0.020640   0.689  0.491
## hypertension -0.007678  0.015995  -0.480  0.631
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##             edf Ref.df      F p-value
## s(age)      9.908e-01     9 13.733 <2e-16 ***
## s(SBP)      6.175e-07     9  0.000  0.765
## s(LDL)      6.648e-07     9  0.000  0.639
## s(bmi)      4.179e+00     9 41.897 <2e-16 ***
## s(time)     7.892e+00     9 44.960 <2e-16 ***
## s(height)   1.234e+00     9  0.278  0.121
## s(weight)   2.262e-06     9  0.000  0.666
## ---

```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.22   Deviance explained = 22.4%
## GCV = 0.27867   Scale est. = 0.27738   n = 5000

# test error
gam_pred = predict(gam_mod, newdata = x_test)
gam_rmse = sqrt(mean((y_test - gam_pred)^2))

cat("Test error for the linear model: ", gam_rmse, "\n")

## Test error for the linear model: 0.5700836
```


Model validation

```
set.seed(1)
resamp = resamples(list(lm = linear_mod,
                       enet =enet_mod,
                        pls = pls_mod,
                        MARS = mars_mod,
                        GAM = gam_mod))

summary(resamp)
```

```
##
## Call:
## summary.resamples(object = resamp)
##
## Models: lm, enet, pls, MARS, GAM
## Number of resamples: 10
##
## MAE
##           Min.    1st Qu.    Median      Mean   3rd Qu.      Max. NA's
## lm    0.4206892 0.4321712 0.4404316 0.4391934 0.4453954 0.4598289    0
## enet  0.4283094 0.4374741 0.4398703 0.4416173 0.4449426 0.4658299    0
## pls   0.4205572 0.4324090 0.4404365 0.4392051 0.4454607 0.4596797    0
## MARS  0.4033817 0.4179151 0.4234065 0.4221856 0.4275553 0.4418075    0
## GAM   0.4042397 0.4194058 0.4242789 0.4228546 0.4275021 0.4421737    0
##
## RMSE
##           Min.    1st Qu.    Median      Mean   3rd Qu.      Max. NA's
## lm    0.5374189 0.5385233 0.5520867 0.5511050 0.5589697 0.5715949    0
## enet  0.5418907 0.5460887 0.5513664 0.5541943 0.5592095 0.5774083    0
## pls   0.5372125 0.5386023 0.5521459 0.5510768 0.5588196 0.5714301    0
## MARS  0.5086158 0.5180140 0.5277205 0.5277239 0.5330971 0.5523110    0
## GAM   0.5104739 0.5194562 0.5286409 0.5286279 0.5334558 0.5525356    0
##
## Rsquared
##           Min.    1st Qu.    Median      Mean   3rd Qu.      Max. NA's
## lm    0.08863443 0.1400413 0.1491380 0.1475987 0.1614569 0.1965523    0
## enet  0.09065183 0.1294793 0.1409665 0.1404523 0.1642418 0.1736794    0
## pls   0.08861598 0.1403744 0.1493405 0.1476804 0.1609204 0.1975762    0
## MARS  0.16452631 0.2016386 0.2149715 0.2179857 0.2395080 0.2656206    0
## GAM   0.16168005 0.1988396 0.2126428 0.2152374 0.2350941 0.2640126    0
```

```
bwplot(resamp, metric = "RMSE")
```



According to the model comparison via 10-fold cross validation, the mean RMSE of the MARS model is the lowest for this specific dataset and predefined seed. Hence, MARS is chosen as the final model to predict the relationship between log-transformed antibody levels and the predictors of interest. Also, MARS offers better interpretability.